

IEOR142_Fiinal_Project_Notebook

May 5, 2023

1 IEOB 142 Final Project - Predicting Spotify Track Popularity

Mina Baghai, Abhigyan Biswas, Toufiq Hossain, Naya Lee, Luna Ragot

2 Data Cleaning

```
[ ]: # loading the dataset
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[1]: import numpy as np
import pandas as pd
```

```
[ ]: #artists = pd.read_csv('/content/drive/MyDrive/IEOB_142_Project/Data_Sets/
↳artists.csv')
#tracks = pd.read_csv('/content/drive/MyDrive/IEOB_142_Project/Data_Sets/tracks.
↳csv')

artists = pd.read_csv('/content/drive/MyDrive/artists.csv')
tracks = pd.read_csv('/content/drive/MyDrive/tracks.csv')
```

```
[3]: artists.head()
```

```
[3]:
```

	id	followers	genres	\
0	ODheY5irMjBUeLybbCUEZ2	0.0	[]	
1	ODlhY15l3wsrnlfGio2bjU	5.0	[]	
2	ODmRESX2JknGPQy015yXg7	0.0	[]	
3	ODmhnbHjm1qw6NCYPeZNgJ	0.0	[]	
4	ODn11fWM7vHQ3rinvWE14E	2.0	[]	

	name	popularity
0	Armid & Amir Zare Pashai feat. Sara Rouzbehani	0
1		0
2	Sadaa	0

3	Tra'gruda	0
4	Ioannis Panoutsopoulos	0

```
[4]: tracks.head()
```

```
[4]:
```

	id	name	popularity	\
0	35iwgR4jXetI318WEWsa1Q	Carve	6	
1	021ht4sdgPcrDgSk7JTbKY	Capítulo 2.16 - Banquero Anarquista	0	
2	07A5yehtSnoedViJAZkNnc	Vivo para Quererte - Remasterizado	0	
3	08FmqUhxyLTn6pAh6bk45	El Prisionero - Remasterizado	0	
4	08y9GfoqCWf0GsKdwojr5e	Lady of the Evening	0	

	duration_ms	explicit	artists	id_artists	\
0	126903	0	['Uli']	['45tIt06XoIOIio4LBEVpls']	
1	98200	0	['Fernando Pessoa']	['14jtPC0oNZwquk5wd9DxrY']	
2	181640	0	['Ignacio Corsini']	['5Li0oJbxVSAMkBS2fUm3X2']	
3	176907	0	['Ignacio Corsini']	['5Li0oJbxVSAMkBS2fUm3X2']	
4	163080	0	['Dick Haymes']	['3BiJGZsyX9sJchTqcSA7Su']	

	release_date	danceability	energy	key	loudness	mode	speechiness	\
0	1922-02-22	0.645	0.4450	0	-13.338	1	0.4510	
1	1922-06-01	0.695	0.2630	0	-22.136	1	0.9570	
2	1922-03-21	0.434	0.1770	1	-21.180	1	0.0512	
3	1922-03-21	0.321	0.0946	7	-27.961	1	0.0504	
4	1922	0.402	0.1580	3	-16.900	0	0.0390	

	acousticness	instrumentalness	liveness	valence	tempo	time_signature
0	0.674	0.7440	0.151	0.127	104.851	3
1	0.797	0.0000	0.148	0.655	102.009	1
2	0.994	0.0218	0.212	0.457	130.418	5
3	0.995	0.9180	0.104	0.397	169.980	3
4	0.989	0.1300	0.311	0.196	103.220	4

```
[5]: #add length of id to see if there's more than one artist
tracks['len artist id'] = [len(i) for i in tracks['id_artists']]

#drop tracks with more than one artist
tracks = tracks[tracks['len artist id'] <= 26]

tracks.head()
```

```
[5]:
```

	id	name	popularity	\
0	35iwgR4jXetI318WEWsa1Q	Carve	6	
1	021ht4sdgPcrDgSk7JTbKY	Capítulo 2.16 - Banquero Anarquista	0	
2	07A5yehtSnoedViJAZkNnc	Vivo para Quererte - Remasterizado	0	
3	08FmqUhxyLTn6pAh6bk45	El Prisionero - Remasterizado	0	
4	08y9GfoqCWf0GsKdwojr5e	Lady of the Evening	0	

	duration_ms	explicit	artists	id_artists	\
0	126903	0	['Uli']	['45tIt06XoIOIio4LBEVpls']	
1	98200	0	['Fernando Pessoa']	['14jtPC0oNZwquk5wd9DxrY']	
2	181640	0	['Ignacio Corsini']	['5Li0oJbxVSAMkBS2fUm3X2']	
3	176907	0	['Ignacio Corsini']	['5Li0oJbxVSAMkBS2fUm3X2']	
4	163080	0	['Dick Haymes']	['3BiJGZsyX9sJchTqcSA7Su']	

	release_date	danceability	energy	...	loudness	mode	speechiness	\
0	1922-02-22	0.645	0.4450	...	-13.338	1	0.4510	
1	1922-06-01	0.695	0.2630	...	-22.136	1	0.9570	
2	1922-03-21	0.434	0.1770	...	-21.180	1	0.0512	
3	1922-03-21	0.321	0.0946	...	-27.961	1	0.0504	
4	1922	0.402	0.1580	...	-16.900	0	0.0390	

	acousticness	instrumentalness	liveness	valence	tempo	time_signature	\
0	0.674	0.7440	0.151	0.127	104.851	3	
1	0.797	0.0000	0.148	0.655	102.009	1	
2	0.994	0.0218	0.212	0.457	130.418	5	
3	0.995	0.9180	0.104	0.397	169.980	3	
4	0.989	0.1300	0.311	0.196	103.220	4	

	len	artist	id
0			26
1			26
2			26
3			26
4			26

[5 rows x 21 columns]

```
[6]: #cleaning id_artists to match id of artists dataset
tracks['simplified_artist_id'] = [i.replace("[", '').replace("]", '') for i_
    ↪in tracks['id_artists']]

#merging datasets
tracks = tracks.merge(artists, left_on='simplified_artist_id', right_on='id')
tracks.head()
```

	id_x	name_x	popularity_x	duration_ms	explicit	\
0	35iwgR4jXetI318WEWsa1Q	Carve	6	126903	0	
1	OPH9AACae1f957JAavh0l2	Lazy Boi	0	157333	0	
2	2SiNuAZ6jIU9xhClRKXcST	Sketch	0	87040	0	
3	4vV7uBcF2AnjNT0ejBS5oL	L'enfer	0	40000	0	
4	598LlBn6jpEpVbLjmZPsYV	Graphite	0	104400	0	

artists	id_artists	release_date	danceability	energy	\
---------	------------	--------------	--------------	--------	---

0	['Uli']	['45tIt06XoIOIio4LBEVpls']	1922-02-22	0.645	0.44500
1	['Uli']	['45tIt06XoIOIio4LBEVpls']	1922-02-22	0.298	0.46000
2	['Uli']	['45tIt06XoIOIio4LBEVpls']	1922-02-22	0.634	0.00399
3	['Uli']	['45tIt06XoIOIio4LBEVpls']	1922-02-22	0.657	0.32500
4	['Uli']	['45tIt06XoIOIio4LBEVpls']	1922-02-22	0.644	0.68400

	...	valence	tempo	time_signature	len	artist	id	\
0	...	0.127	104.851	3			26	
1	...	0.402	87.921	4			26	
2	...	0.396	79.895	4			26	
3	...	0.105	81.944	5			26	
4	...	0.138	100.031	4			26	

	simplified_artist_id	id_y	followers	genres	name_y	\
0	45tIt06XoIOIio4LBEVpls	45tIt06XoIOIio4LBEVpls	91.0	[]	Uli	
1	45tIt06XoIOIio4LBEVpls	45tIt06XoIOIio4LBEVpls	91.0	[]	Uli	
2	45tIt06XoIOIio4LBEVpls	45tIt06XoIOIio4LBEVpls	91.0	[]	Uli	
3	45tIt06XoIOIio4LBEVpls	45tIt06XoIOIio4LBEVpls	91.0	[]	Uli	
4	45tIt06XoIOIio4LBEVpls	45tIt06XoIOIio4LBEVpls	91.0	[]	Uli	

	popularity_y
0	4
1	4
2	4
3	4
4	4

[5 rows x 27 columns]

```
[7]: #cleaning columns and column names
tracks.rename(columns={'id_x': 'track_id', 'name_x': 'track_name',
    → 'popularity_x': 'track_popularity', 'simplified_artist_id': 'artist_id',
    → 'name_y': 'artist_name', 'popularity_y': 'artist_popularity'}, inplace=True)
tracks.drop(['artists', 'id_artists', 'id_y', 'len artist id'], axis=1,
    → inplace=True)
tracks.head()
```

[7]:	track_id	track_name	track_popularity	duration_ms	explicit	\
0	35iwgR4jXetI318WEWsa1Q	Carve	6	126903	0	
1	OPH9AACae1f957JAavh0l2	Lazy Boi	0	157333	0	
2	2SiNuAZ6jIU9xhClRKXcST	Sketch	0	87040	0	
3	4vV7uBcF2AnjNT0ejBS5oL	L'enfer	0	40000	0	
4	598LlBn6jpEpVbLjmZPsYV	Graphite	0	104400	0	

	release_date	danceability	energy	key	loudness	...	instrumentalness	\
0	1922-02-22	0.645	0.44500	0	-13.338	...	0.744	
1	1922-02-22	0.298	0.46000	1	-18.645	...	0.856	

2	1922-02-22	0.634	0.00399	5	-29.973	...	0.919
3	1922-02-22	0.657	0.32500	10	-14.319	...	0.856
4	1922-02-22	0.644	0.68400	7	-8.247	...	0.802

	liveness	valence	tempo	time_signature	artist_id	\
0	0.1510	0.127	104.851	3	45tIt06XoIOIio4LBEVpls	
1	0.4360	0.402	87.921	4	45tIt06XoIOIio4LBEVpls	
2	0.1050	0.396	79.895	4	45tIt06XoIOIio4LBEVpls	
3	0.0931	0.105	81.944	5	45tIt06XoIOIio4LBEVpls	
4	0.0847	0.138	100.031	4	45tIt06XoIOIio4LBEVpls	

	followers	genres	artist_name	artist_popularity
0	91.0	[]	Uli	4
1	91.0	[]	Uli	4
2	91.0	[]	Uli	4
3	91.0	[]	Uli	4
4	91.0	[]	Uli	4

[5 rows x 23 columns]

```
[8]: #change date column from objects to int objects for modeling
#tracks['release_date'] = pd.to_datetime(tracks['release_date'], format = "%Y-%m-%d")
tracks['release_date'] = pd.to_datetime(tracks["release_date"]).dt.
    strftime("%Y%m%d")
```

Cleaning up genres column

```
[9]: tracks.set_index('track_id', inplace=True)
```

```
[10]: #changing from string to list of strings
clean = [i.replace("[", '').replace("]", '') for i in tracks['genres']]
clean = [k[1:-1].split("'", '"') for k in clean]

tracks['clean_genres'] = clean
tracks.drop('genres', inplace = True, axis=1)
```

```
[11]: #onehotencoding genres
from sklearn.preprocessing import MultiLabelBinarizer

mlb = MultiLabelBinarizer(sparse_output=True)

df = tracks.join(
    pd.DataFrame.sparse.from_spmatrix(
        mlb.fit_transform(tracks.pop('clean_genres')),
        index=tracks.index,
        columns=mlb.classes_))
```

```
[12]: #only keeping genres that appear more than 8000 times in the dataset
```

```
drops = []
for i in range(4521 - 21):
    if df.iloc[:, i+21].sum() < 8000:
        drops.append(df.iloc[:, i+21].name)

df.drop(drops, axis=1, inplace=True)
tracks = df
```

```
[13]: #replace all spaces in column names with underscores to make modeling easier
```

```
tracks.columns = tracks.columns.str.replace(' ', '_')
tracks.head()
```

```
[13]:
```

	track_name	track_popularity	duration_ms	explicit	\
track_id					
35iwgR4jXetI318WEWsa1Q	Carve	6	126903	0	
OPH9AACae1f957JAavh012	Lazy Boi	0	157333	0	
2SiNuAZ6jIU9xhClRKXcST	Sketch	0	87040	0	
4vV7uBcF2AnjNT0ejBS5oL	L'enfer	0	40000	0	
598LlBn6jpEpVbLjmZPsYV	Graphite	0	104400	0	

	release_date	danceability	energy	key	loudness	\
track_id						
35iwgR4jXetI318WEWsa1Q	19220222	0.645	0.44500	0	-13.338	
OPH9AACae1f957JAavh012	19220222	0.298	0.46000	1	-18.645	
2SiNuAZ6jIU9xhClRKXcST	19220222	0.634	0.00399	5	-29.973	
4vV7uBcF2AnjNT0ejBS5oL	19220222	0.657	0.32500	10	-14.319	
598LlBn6jpEpVbLjmZPsYV	19220222	0.644	0.68400	7	-8.247	

	mode	...	latin	latin_pop	lounge	mellow_gold	\
track_id		...					
35iwgR4jXetI318WEWsa1Q	1	...	0	0	0	0	
OPH9AACae1f957JAavh012	1	...	0	0	0	0	
2SiNuAZ6jIU9xhClRKXcST	0	...	0	0	0	0	
4vV7uBcF2AnjNT0ejBS5oL	0	...	0	0	0	0	
598LlBn6jpEpVbLjmZPsYV	1	...	0	0	0	0	

	psychedelic_rock	rock	rock_en_espanol	soft_rock	\
track_id					
35iwgR4jXetI318WEWsa1Q	0	0	0	0	
OPH9AACae1f957JAavh012	0	0	0	0	
2SiNuAZ6jIU9xhClRKXcST	0	0	0	0	
4vV7uBcF2AnjNT0ejBS5oL	0	0	0	0	
598LlBn6jpEpVbLjmZPsYV	0	0	0	0	

	soul	vocal_jazz
track_id		

35iwgR4jXetI318WEWsa1Q	0	0
0PH9AACae1f957JAavh0l2	0	0
2SiNuAZ6jIU9xhClRKXcST	0	0
4vV7uBcF2AnjNT0ejBS5oL	0	0
598LlBn6jpEpVbLjmZPsYV	0	0

[5 rows x 46 columns]

3 Regression Modeling

Splitting into training and testing set. Same training/testing datasets will be used throughout the regression and classification modeling sections.

```
[14]: from sklearn.model_selection import train_test_split

train, test = train_test_split(tracks, test_size=0.3, random_state=88)

X_train = train.drop('track_popularity', axis=1)
y_train = train['track_popularity']
X_test = test.drop('track_popularity', axis=1)
y_test = test['track_popularity']
print(X_train.shape, X_test.shape)
```

(329026, 45) (141012, 45)

3.1 OLS

```
[15]: #calculate VIFs
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm

def VIF(df, cols):
    values = sm.add_constant(df[cols]).values
    vif = [variance_inflation_factor(values, i) for i in range(len(cols)+1)]
    return pd.Series(vif[1:], index=cols)

VIF(train,
    → ['duration_ms', 'explicit', 'danceability', 'energy', 'key', 'loudness', 'mode', 'acousticness',
        → 'instrumentalness', 'liveness', 'valence', 'tempo', 'time_signature', 'followers', 'artist_popularity'])
```

```
[15]: duration_ms      1.055264
explicit      1.051991
danceability   1.537024
energy        4.149880
key           1.018456
loudness      2.502978
```

```

mode                1.027044
acousticness        2.165865
instrumentalness     1.134507
liveness            1.078624
valence             1.705608
tempo               1.100166
time_signature      1.051884
followers           1.228751
artist_popularity   1.303838
dtype: float64

```

No need to drop any features due to high VIFs since they are all below 5

```

[16]: import statsmodels.formula.api as smf

ols = smf.ols(formula='track_popularity ~ duration_ms + explicit + danceability_
↳ energy + key + loudness + mode + acousticness + instrumentalness +_
↳ liveness + valence + tempo + time_signature + followers + artist_popularity',
              data=train).fit()

print(ols.summary())

```

OLS Regression Results

```

=====
Dep. Variable:      track_popularity    R-squared:                0.423
Model:              OLS                Adj. R-squared:         0.423
Method:             Least Squares       F-statistic:           1.610e+04
Date:               Fri, 05 May 2023    Prob (F-statistic):      0.00
Time:               14:13:57            Log-Likelihood:        -1.3180e+06
No. Observations:   329026              AIC:                  2.636e+06
Df Residuals:       329010              BIC:                  2.636e+06
Df Model:           15
Covariance Type:    nonrobust
=====

```

```

=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
Intercept      8.9785      0.319     28.130      0.000      8.353
9.604
duration_ms    6.127e-06    2.06e-07    29.723      0.000     5.72e-06
6.53e-06
explicit       6.6870      0.129     51.799      0.000      6.434
6.940
danceability   12.9013      0.180     71.481      0.000     12.548
13.255
energy        -0.7028      0.193     -3.635      0.000     -1.082

```


-0.324					
key	0.0252	0.007	3.800	0.000	0.012
0.038					
loudness	0.5317	0.008	67.067	0.000	0.516
0.547					
mode	-0.1392	0.050	-2.799	0.005	-0.237
-0.042					
acousticness	-7.7630	0.101	-77.229	0.000	-7.960
-7.566					
instrumentalness	-3.9817	0.102	-39.030	0.000	-4.182
-3.782					
liveness	-6.4951	0.130	-50.113	0.000	-6.749
-6.241					
valence	-8.6830	0.120	-72.491	0.000	-8.918
-8.448					
tempo	0.0167	0.001	20.401	0.000	0.015
0.018					
time_signature	0.3926	0.052	7.524	0.000	0.290
0.495					
followers	-1.28e-07	6.68e-09	-19.153	0.000	-1.41e-07
-1.15e-07					
artist_popularity	0.4513	0.001	335.872	0.000	0.449
0.454					
=====					
Omnibus:	276.956	Durbin-Watson:	2.009		
Prob(Omnibus):	0.000	Jarque-Bera (JB):	308.399		
Skew:	0.031	Prob(JB):	1.08e-67		
Kurtosis:	3.137	Cond. No.	5.98e+07		
=====					

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.98e+07. This might indicate that there are strong multicollinearity or other numerical problems.

All of the p-values are now extremely small, so we will not be removing any more features.

```
[ ]: # compute out-of-sample R-squared using the test set
def OSR2(model, df_train, df_test, dependent_var):
    y_test = df_test[dependent_var]
    y_pred = model.predict(df_test)
    SSE = np.sum((y_test - y_pred)**2)
    SST = np.sum((y_test - np.mean(df_train[dependent_var]))**2)
    return 1 - SSE/SST
```

```
[ ]: # compute test set RSS
def test_rss(model, df_train, df_test, dependent_var):
```

```

y_test = df_test[dependent_var]
y_pred = model.predict(df_test)
return np.sum((y_test - y_pred)**2)

```

```

[ ]: ols_osr2 = OSR2(ols, train, test, 'track_popularity')
ols_test_rss = test_rss(ols, train, test, 'track_popularity')
print('Out-of-sample R-squared for OLS: ', ols_osr2)
print('Test RSS for OLS: ', ols_test_rss)

```

Out-of-sample R-squared for OLS: 0.418832640795064
Test RSS for OLS: 25052477.51463911

3.2 CART - Regression Tree

```

[ ]: import matplotlib.pyplot as plt
from sklearn.tree import plot_tree
from sklearn.tree import DecisionTreeRegressor, DecisionTreeClassifier
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
from sklearn.model_selection import KFold

```

```

[ ]: dtr_X_train = X_train.drop(['track_name', 'artist_name', 'artist_id'], axis=1)
dtr_X_test = X_test.drop(['track_name', 'artist_name', 'artist_id'], axis=1)

dtr = DecisionTreeRegressor(min_samples_split=10,
                             ccp_alpha=0.02,
                             random_state = 88)
dtr.fit(dtr_X_train, y_train)

```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:768:
UserWarning: pandas.DataFrame with sparse columns found.It will be converted to
a dense numpy array.
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/utils/_array_api.py:185:
FutureWarning: The behavior of .astype from SparseDtype to a non-sparse dtype is
deprecated. In a future version, this will return a non-sparse array with the
requested dtype. To retain the old behavior, use
`obj.astype(SparseDtype(dtype))`
array = numpy.asarray(array, order=order, dtype=dtype)

```

[ ]: DecisionTreeRegressor(ccp_alpha=0.02, min_samples_split=10, random_state=88)

```

```

[ ]: print('Node count =', dtr.tree_.node_count)
plt.figure(figsize=(9,7))
plot_tree(dtr,
           feature_names=dtr_X_train.columns,
           class_names=['0', '1'],
           filled=True,


```

```

        impurity=False,
        rounded=True,
        fontsize=12)
plt.show()

```

Node count = 843



```

[ ]: def OSR2(model, X_test, y_test, y_train):
    y_pred = model.predict(X_test)
    SSE = np.sum((y_test - y_pred)**2)
    SST = np.sum((y_test - np.mean(y_train))**2)
    return (1 - SSE/SST)

```

```

def test_RSS(model, X_test, y_test, y_train):
    y_pred = model.predict(X_test)
    return np.sum((y_test - y_pred)**2)

```

```

[ ]: dtr_osr2 = OSR2(dtr, dtr_X_test, y_test, y_train)
    dtr_test_rss = test_RSS(dtr, dtr_X_test, y_test, y_train)
    print('Out-of-sample R-squared for Regression Tree:', dtr_osr2)
    print('Test RSS for Regression Tree: ', dtr_test_rss)

```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:768:
UserWarning: pandas.DataFrame with sparse columns found.It will be converted to
a dense numpy array.

```
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/utils/_array_api.py:185:
FutureWarning: The behavior of .astype from SparseDtype to a non-sparse dtype is
deprecated. In a future version, this will return a non-sparse array with the
requested dtype. To retain the old behavior, use
`obj.astype(SparseDtype(dtype))`
    array = numpy.asarray(array, order=order, dtype=dtype)
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:768:
UserWarning: pandas.DataFrame with sparse columns found.It will be converted to
a dense numpy array.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/utils/_array_api.py:185:
FutureWarning: The behavior of .astype from SparseDtype to a non-sparse dtype is
deprecated. In a future version, this will return a non-sparse array with the
requested dtype. To retain the old behavior, use
`obj.astype(SparseDtype(dtype))`
    array = numpy.asarray(array, order=order, dtype=dtype)

Out-of-sample R-squared for Regression Tree: 0.6334125297223259
Test RSS for Regression Tree: 15802546.73773128
```

3.3 Comparison of Regression Models (OLS vs. Regression Tree)


```
[ ]: import matplotlib.pyplot as plt

models = ['OLS', 'Regression Tree']
osr2s = [ols_osr2, dtr_osr2]
test_RSSs = [ols_test_rss, dtr_test_rss]

plt.figure(figsize=(15,4))

plt.subplot(1,2,1)
plt.bar(models, osr2s)
plt.ylabel('Out of Sample R2 Scores')
plt.title('Out of Sample R2 Score Model Comparison')

plt.subplot(1,2,2)
plt.bar(models, test_RSSs)
plt.ylabel('Testing RSS Scores')
plt.title('Testing RSS Score Model Comparison');
```




Comparing our two regression models, we can clearly see that the regression tree model performs much better compared to the OLS model. The regression tree model has a significantly higher OSR2 score at above 0.6 compared to the OLS model having an OSR2 score much lower around 0.4. Similarly, when looking at the testing RSS scores, the regression tree has a score of around 9000000 lower than that of the OLS model.

4 Classification Modeling

Considering tracks with a track popularity over 40 to be considered popular. Classification models to predict whether a track is popular or not.

4.1 Logistic Regression

```
[ ]: plt.hist(train['track_popularity'])
      #Probably will make a break along the lines of 40 track popularity for 1/0 for
      ↳ logistic regression model
      #since it's a 50/50 ish split
      tracks['over40'] = (tracks['track_popularity'] > 40).astype(int)
```



```
[ ]: train, test = train_test_split(tracks, test_size=0.3, random_state=88)
X_train = train.drop('track_popularity', axis=1)
y_train = train['track_popularity']
X_test = test.drop('track_popularity', axis=1)
y_test = test['track_popularity']
print(X_train.shape, X_test.shape)
```

```
(329026, 46) (141012, 46)
```

```
[ ]: VIF(train,
↳ ['duration_ms', 'explicit', 'danceability', 'energy', 'key', 'loudness', 'mode', 'acousticness',
↳
↳ 'instrumentalness', 'liveness', 'valence', 'tempo', 'time_signature', 'followers', 'artist_popularity'])
```

```
[ ]: duration_ms      1.055264
explicit            1.051991
danceability        1.537024
energy              4.149880
key                 1.018456
loudness            2.502978
mode                1.027044
acousticness        2.165865
```

```
instrumentalness    1.134507
liveness            1.078624
valence             1.705608
tempo              1.100166
time_signature      1.051884
followers           1.228751
artist_popularity   1.303838
dtype: float64
```

```
[ ]: VIF(train,
↳ ['duration_ms', 'explicit', 'danceability', 'key', 'loudness', 'mode', 'acousticness',
↳
↳ 'instrumentalness', 'liveness', 'valence', 'tempo', 'time_signature', 'followers', 'artist_popularity'])
```

```
[ ]: duration_ms      1.054485
explicit             1.051468
danceability         1.518525
key                  1.018353
loudness             1.421752
mode                 1.024924
acousticness         1.412732
instrumentalness     1.107882
liveness             1.028398
valence              1.515340
tempo                1.097532
time_signature       1.049389
followers            1.227288
artist_popularity    1.301911
dtype: float64
```

Removed energy since the VIF was close to 5. Now all VIFs are below 2. This seems reasonable to have for my model. We will cut down more columns if they are statistically not significant or otherwise don't help our model.

```
[ ]: train.columns
```

```
[ ]: Index(['track_name', 'track_popularity', 'duration_ms', 'explicit',
'release_date', 'danceability', 'energy', 'key', 'loudness', 'mode',
'speechiness', 'acousticness', 'instrumentalness', 'liveness',
'valence', 'tempo', 'time_signature', 'artist_id', 'followers',
'artist_name', 'artist_popularity', '', 'adult_standards', 'album_rock',
'art_rock', 'brill_building_pop', 'c-pop', 'classic_rock', 'cool_jazz',
'country_rock', 'filmi', 'folk', 'folk_rock', 'hard_rock', 'hoerspiel',
'jazz', 'latin', 'latin_pop', 'lounge', 'mellow_gold',
'psychedelic_rock', 'rock', 'rock_en_espanol', 'soft_rock', 'soul',
'vocal_jazz', 'over40'],
dtype='object')
```

```
[ ]: logit = smf.logit(formula='over40 ~ duration_ms + explicit + danceability + key_
    ↳+ loudness + mode + speechiness + acousticness + \
        instrumentalness + liveness + valence + tempo +_
    ↳time_signature + followers',
        data=train).fit()

print(logit.summary())
```

Optimization terminated successfully.

Current function value: 0.491185

Iterations 7

Logit Regression Results

```
=====
Dep. Variable:          over40    No. Observations:          329026
Model:                  Logit     Df Residuals:            329011
Method:                  MLE      Df Model:                14
Date:                   Fri, 05 May 2023    Pseudo R-squ.:          0.1284
Time:                   17:35:22    Log-Likelihood:         -1.6161e+05
converged:              True      LL-Null:                -1.8542e+05
Covariance Type:        nonrobust    LLR p-value:            0.000
=====
```

```
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
----
Intercept      -0.7284      0.059     -12.421     0.000     -0.843
-0.613
duration_ms    2.611e-07    4.09e-08      6.382     0.000     1.81e-07
3.41e-07
explicit        1.1424      0.022     52.572     0.000      1.100
1.185
danceability    2.2753      0.035     65.466     0.000      2.207
2.343
key             0.0053      0.001      4.288     0.000      0.003
0.008
loudness        0.1028      0.001     77.821     0.000      0.100
0.105
mode            0.0291      0.009      3.123     0.002      0.011
0.047
speechiness     -1.3848      0.041    -33.375     0.000     -1.466
-1.303
acousticness    -0.7134      0.016    -45.620     0.000     -0.744
-0.683
instrumentalness -0.5842      0.023    -25.728     0.000     -0.629
-0.540
liveness        -0.5521      0.026    -21.246     0.000     -0.603
=====
```


-0.501					
valence	-1.5141	0.021	-72.054	0.000	-1.555
-1.473					
tempo	0.0024	0.000	14.758	0.000	0.002
0.003					
time_signature	0.0252	0.012	2.103	0.036	0.002
0.049					
followers	1.166e-07	1.68e-09	69.350	0.000	1.13e-07
1.2e-07					

=====

====

Adding genres into the model

```
[ ]: logit2 = smf.logit(formula='over40 ~ duration_ms + explicit + danceability +
    ↳key + loudness + mode + speechiness + acousticness \
        + instrumentalness + liveness + valence + tempo +
    ↳time_signature + followers \
        + adult_standards + album_rock + art_rock +
    ↳brill_building_pop + classic_rock + cool_jazz \
        + country_rock + filmi + folk + folk_rock + hard_rock
    ↳+ hoerspiel + jazz + latin + latin_pop + lounge + mellow_gold \
        + psychedelic_rock + rock + rock_en_espanol +
    ↳soft_rock + soul + vocal_jazz',
        data=train).fit()

print(logit2.summary())
```

Optimization terminated successfully.

Current function value: 0.484597

Iterations 7

Logit Regression Results

Dep. Variable:	over40	No. Observations:	329026
Model:	Logit	Df Residuals:	328988
Method:	MLE	Df Model:	37
Date:	Fri, 05 May 2023	Pseudo R-squ.:	0.1401
Time:	17:35:27	Log-Likelihood:	-1.5944e+05
converged:	True	LL-Null:	-1.8542e+05
Covariance Type:	nonrobust	LLR p-value:	0.000

=====

=====

	coef	std err	z	P> z	[0.025
--	------	---------	---	------	--------

0.975]

Intercept	-0.7530	0.059	-12.661	0.000	-0.870
	-0.636				

duration_ms	2.571e-07	4.2e-08	6.119	0.000	1.75e-07
3.39e-07					
explicit	1.1274	0.022	51.470	0.000	1.084
1.170					
danceability	2.2922	0.036	63.956	0.000	2.222
2.362					
key	0.0049	0.001	3.914	0.000	0.002
0.007					
loudness	0.1047	0.001	76.990	0.000	0.102
0.107					
mode	0.0054	0.009	0.572	0.567	-0.013
0.024					
speechiness	-0.9385	0.047	-20.125	0.000	-1.030
-0.847					
acousticness	-0.5714	0.016	-34.915	0.000	-0.604
-0.539					
instrumentalness	-0.5180	0.023	-22.441	0.000	-0.563
-0.473					
liveness	-0.5564	0.026	-21.291	0.000	-0.608
-0.505					
valence	-1.4981	0.021	-70.175	0.000	-1.540
-1.456					
tempo	0.0022	0.000	13.513	0.000	0.002
0.002					
time_signature	0.0092	0.012	0.759	0.448	-0.015
0.033					
followers	1.049e-07	1.74e-09	60.343	0.000	1.02e-07
1.08e-07					
adult_standards	0.0386	0.036	1.085	0.278	-0.031
0.108					
album_rock	0.1072	0.039	2.749	0.006	0.031
0.184					
art_rock	-0.0484	0.032	-1.515	0.130	-0.111
0.014					
brill_building_pop	-0.1925	0.043	-4.487	0.000	-0.277
-0.108					
classic_rock	-0.9658	0.036	-26.524	0.000	-1.037
-0.894					
cool_jazz	-0.5566	0.112	-4.976	0.000	-0.776
-0.337					
country_rock	0.5084	0.034	14.885	0.000	0.441
0.575					
filmi	-0.8134	0.046	-17.859	0.000	-0.903
-0.724					
folk	0.0749	0.044	1.717	0.086	-0.011
0.160					
folk_rock	-0.0675	0.041	-1.639	0.101	-0.148
0.013					

hard_rock	-0.3215	0.037	-8.765	0.000	-0.393
-0.250					
hoerspiel	-0.9276	0.072	-12.832	0.000	-1.069
-0.786					
jazz	-0.0918	0.106	-0.867	0.386	-0.299
0.116					
latin	0.1027	0.036	2.852	0.004	0.032
0.173					
latin_pop	0.8019	0.039	20.686	0.000	0.726
0.878					
lounge	-0.2878	0.052	-5.527	0.000	-0.390
-0.186					
mellow_gold	0.0598	0.043	1.376	0.169	-0.025
0.145					
psychedelic_rock	-0.1998	0.042	-4.786	0.000	-0.282
-0.118					
rock	0.8905	0.026	33.826	0.000	0.839
0.942					
rock_en_espanol	0.4370	0.028	15.806	0.000	0.383
0.491					
soft_rock	0.4052	0.038	10.739	0.000	0.331
0.479					
soul	0.2393	0.031	7.636	0.000	0.178
0.301					
vocal_jazz	-0.2184	0.053	-4.110	0.000	-0.323
-0.114					

=====

=====

Remove mode, time_signature, adult_standards, art_rock, and jazz because they all have very high p-values.

```
[ ]: logit3 = smf.logit(formula='over40 ~ duration_ms + explicit + danceability +
    ↳key + loudness + speechiness + acousticness \
    + instrumentalness + liveness + valence + tempo +
    ↳followers \
    + album_rock + brill_building_pop + classic_rock +
    ↳cool_jazz \
    + country_rock + filmi + folk + folk_rock + hard_rock,
    ↳+ hoerspiel + latin + latin_pop + lounge + mellow_gold \
    + psychedelic_rock + rock + rock_en_espanol +
    ↳soft_rock + soul + vocal_jazz',
    data=train).fit()

print(logit3.summary())
```

Optimization terminated successfully.

Current function value: 0.484605

Iterations 7

Logit Regression Results

```

=====
Dep. Variable:          over40    No. Observations:          329026
Model:                  Logit      Df Residuals:              328993
Method:                 MLE        Df Model:                  32
Date:                  Fri, 05 May 2023    Pseudo R-squ.:            0.1401
Time:                  17:35:34    Log-Likelihood:          -1.5945e+05
converged:              True        LL-Null:                 -1.8542e+05
Covariance Type:        nonrobust    LLR p-value:             0.000
=====

```

```

=====
                                coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
Intercept                   -0.7120      0.036     -19.793      0.000     -0.782
-0.641
duration_ms                 2.534e-07    4.2e-08      6.038      0.000     1.71e-07
3.36e-07
explicit                     1.1273      0.022     51.491      0.000      1.084
1.170
danceability                 2.2923      0.036     64.394      0.000      2.223
2.362
key                          0.0048      0.001      3.874      0.000      0.002
0.007
loudness                     0.1048      0.001     77.277      0.000      0.102
0.107
speechiness                 -0.9417      0.047    -20.228      0.000     -1.033
-0.850
acousticness                -0.5712      0.016    -35.162      0.000     -0.603
-0.539
instrumentalness            -0.5196      0.023    -22.525      0.000     -0.565
-0.474
liveness                    -0.5557      0.026    -21.270      0.000     -0.607
-0.505
valence                     -1.4969      0.021    -70.176      0.000     -1.539
-1.455
tempo                       0.0022      0.000     13.497      0.000      0.002
0.002
followers                   1.052e-07    1.73e-09     60.726      0.000     1.02e-07
1.09e-07
album_rock                   0.0935      0.038      2.448      0.014      0.019
0.168
brill_building_pop         -0.1753      0.041     -4.316      0.000     -0.255
-0.096
classic_rock                -0.9680      0.036    -26.621      0.000     -1.039
-0.897

```

cool_jazz	-0.6416	0.062	-10.365	0.000	-0.763
-0.520					
country_rock	0.5101	0.034	15.013	0.000	0.444
0.577					
filmi	-0.8153	0.046	-17.906	0.000	-0.905
-0.726					
folk	0.0771	0.044	1.768	0.077	-0.008
0.163					
folk_rock	-0.0711	0.041	-1.728	0.084	-0.152
0.010					
hard_rock	-0.3126	0.036	-8.604	0.000	-0.384
-0.241					
hoerspiel	-0.9293	0.072	-12.861	0.000	-1.071
-0.788					
latin	0.1022	0.036	2.836	0.005	0.032
0.173					
latin_pop	0.8018	0.039	20.684	0.000	0.726
0.878					
lounge	-0.2735	0.049	-5.629	0.000	-0.369
-0.178					
mellow_gold	0.0661	0.043	1.530	0.126	-0.019
0.151					
psychedelic_rock	-0.2144	0.041	-5.247	0.000	-0.294
-0.134					
rock	0.8824	0.026	34.048	0.000	0.832
0.933					
rock_en_espanol	0.4372	0.028	15.812	0.000	0.383
0.491					
soft_rock	0.4077	0.037	10.886	0.000	0.334
0.481					
soul	0.2418	0.031	7.817	0.000	0.181
0.302					
vocal_jazz	-0.2066	0.049	-4.214	0.000	-0.303
-0.111					

=====

=====

Now remove folk, folk_rock, and mellow_gold as they are the remaining features with high p-values.

```
[ ]: logit4 = smf.logit(formula='over40 ~ duration_ms + explicit + danceability +
    ↳key + loudness + speechiness + acousticness \
    + instrumentalness + liveness + valence + tempo +
    ↳followers \
    + album_rock + brill_building_pop + classic_rock +
    ↳cool_jazz \
    + country_rock + filmi + hard_rock + hoerspiel +
    ↳latin + latin_pop + lounge \
```

```

+ psychedelic_rock + rock + rock_en_espanol +
↪soft_rock + soul + vocal_jazz',
    data=train).fit()

print(logit4.summary())

```

Optimization terminated successfully.

Current function value: 0.484615

Iterations 7

Logit Regression Results

```

=====
Dep. Variable:          over40    No. Observations:          329026
Model:                  Logit      Df Residuals:              328996
Method:                  MLE        Df Model:                  29
Date:                   Fri, 05 May 2023    Pseudo R-squ.:            0.1401
Time:                   17:35:38    Log-Likelihood:           -1.5945e+05
converged:              True        LL-Null:                  -1.8542e+05
Covariance Type:        nonrobust    LLR p-value:              0.000
=====
=====

```

	coef	std err	z	P> z	[0.025
0.975]					

Intercept	-0.7127	0.036	-19.815	0.000	-0.783
-0.642					
duration_ms	2.538e-07	4.2e-08	6.048	0.000	1.72e-07
3.36e-07					
explicit	1.1273	0.022	51.493	0.000	1.084
1.170					
danceability	2.2918	0.036	64.382	0.000	2.222
2.362					
key	0.0048	0.001	3.874	0.000	0.002
0.007					
loudness	0.1046	0.001	77.254	0.000	0.102
0.107					
speechiness	-0.9420	0.047	-20.237	0.000	-1.033
-0.851					
acousticness	-0.5703	0.016	-35.142	0.000	-0.602
-0.538					
instrumentalness	-0.5205	0.023	-22.575	0.000	-0.566
-0.475					
liveness	-0.5558	0.026	-21.271	0.000	-0.607
-0.505					
valence	-1.4974	0.021	-70.204	0.000	-1.539
-1.456					
tempo	0.0022	0.000	13.511	0.000	0.002

0.002					
followers	1.054e-07	1.73e-09	61.031	0.000	1.02e-07
1.09e-07					
album_rock	0.1010	0.038	2.689	0.007	0.027
0.175					
brill_building_pop	-0.1779	0.039	-4.560	0.000	-0.254
-0.101					
classic_rock	-0.9690	0.036	-27.224	0.000	-1.039
-0.899					
cool_jazz	-0.6435	0.062	-10.397	0.000	-0.765
-0.522					
country_rock	0.5182	0.031	16.650	0.000	0.457
0.579					
filmi	-0.8161	0.046	-17.925	0.000	-0.905
-0.727					
hard_rock	-0.3237	0.036	-9.014	0.000	-0.394
-0.253					
hoerspiel	-0.9298	0.072	-12.868	0.000	-1.071
-0.788					
latin	0.1019	0.036	2.828	0.005	0.031
0.173					
latin_pop	0.8010	0.039	20.665	0.000	0.725
0.877					
lounge	-0.2744	0.049	-5.654	0.000	-0.369
-0.179					
psychedelic_rock	-0.2341	0.040	-5.893	0.000	-0.312
-0.156					
rock	0.8865	0.026	34.366	0.000	0.836
0.937					
rock_en_espanol	0.4372	0.028	15.813	0.000	0.383
0.491					
soft_rock	0.4483	0.026	17.238	0.000	0.397
0.499					
soul	0.2432	0.031	7.872	0.000	0.183
0.304					
vocal_jazz	-0.2040	0.049	-4.162	0.000	-0.300
-0.108					

```
=====
=====
```

```
[ ]: import numpy as np
      np.exp(2.2918) - 1
```

```
[ ]: 8.892728576001097
```

We can see that the odds of a track being popular goes up by almost 900% when the danceability score goes up by 1.

Now all of our p-values are low, indicating that all of our features are significant and thus we will keep all the remaining features and evaluate our 4th logistic regression model on the testing data.

```
[ ]: from sklearn.metrics import confusion_matrix
```

To remind you of what each element of the confusion matrix represents:

TN FP FN TP

```
[ ]: y_test = test['over40']
y_prob = logit4.predict(X_test)
y_pred = pd.Series([1 if x > 0.35 else 0 for x in y_prob], index=y_prob.index)
cm = confusion_matrix(y_test, y_pred)
print ("Confusion Matrix : \n", cm)
def get_accuracy(cm):
    return (cm.ravel()[0]+cm.ravel()[3])/sum(cm.ravel())
def get_tpr(cm):
    return (cm.ravel()[3]/(cm.ravel()[3] + cm.ravel()[2]))
def get_fpr(cm):
    return (cm.ravel()[1]/(cm.ravel()[1] + cm.ravel()[0]))

logit_accuracy = get_accuracy(cm)
logit_tpr = get_tpr(cm)
logit_fpr = get_fpr(cm)

print('Accuracy: ' + str(logit_accuracy))
print('TPR: ' + str(logit_tpr))
print('FPR: ' + str(logit_fpr))
```

```
Confusion Matrix :
[[90799 14841]
 [19088 16284]]
Accuracy: 0.7593892718350211
TPR: 0.46036412982019675
FPR: 0.1404865581219235
```

4.2 CART - Classification Tree

```
[ ]: dtc_X_train = X_train.drop(['track_name', 'artist_name', 'artist_id',
    ↳ 'over40'], axis=1)
dtc_X_test = X_test.drop(['track_name', 'artist_name', 'artist_id', 'over40'],
    ↳ axis=1)
```

```
[ ]: dtc = DecisionTreeClassifier(min_samples_split=20,
    ccp_alpha=0.00,
    random_state = 88)
dtc.fit(dtc_X_train, y_train)
```

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:768:

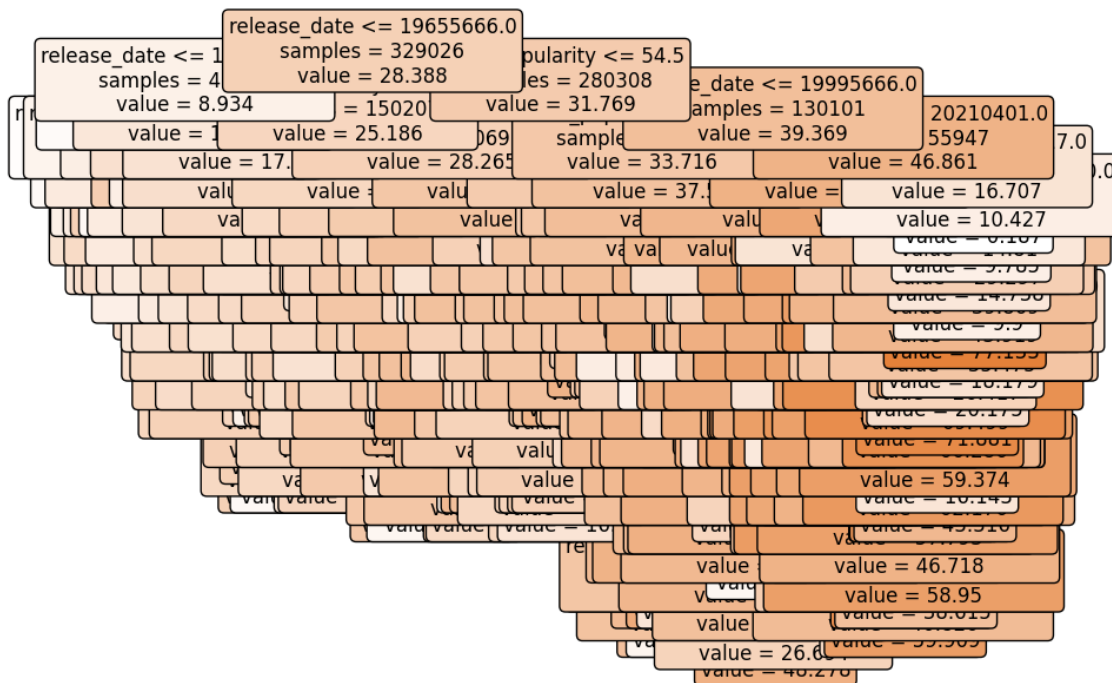
UserWarning: pandas.DataFrame with sparse columns found.It will be converted to a dense numpy array.

```
warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/utils/_array_api.py:185:
FutureWarning: The behavior of .astype from SparseDtype to a non-sparse dtype is
deprecated. In a future version, this will return a non-sparse array with the
requested dtype. To retain the old behavior, use
`obj.astype(SparseDtype(dtype))`
array = numpy.asarray(array, order=order, dtype=dtype)
```

```
[ ]: DecisionTreeClassifier(min_samples_split=20, random_state=88)
```

```
[ ]: print('Node count =', dtr.tree_.node_count)
plt.figure(figsize=(9,7))
plot_tree(dtr,
          feature_names=dtr_X_train.columns,
          class_names=['0', '1'],
          filled=True,
          impurity=False,
          rounded=True,
          fontsize=12)
plt.show()
```

Node count = 843



```
[ ]: y_pred_ct = dtc.predict(dtc_X_test)
y_pred_ct_thresh = [1 if x >= 38 else 0 for x in y_pred_ct]
cm_ct = confusion_matrix(y_test, y_pred_ct_thresh)
ct_accuracy = get_accuracy(cm_ct)
ct_tpr = get_tpr(cm_ct)
ct_fpr = get_fpr(cm_ct)

print('Accuracy: ', ct_accuracy)
print('TPR: ', ct_tpr)
print('FPR: ', ct_fpr)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:768:
UserWarning: pandas.DataFrame with sparse columns found.It will be converted to
a dense numpy array.
    warnings.warn(
/usr/local/lib/python3.10/dist-packages/sklearn/utils/_array_api.py:185:
FutureWarning: The behavior of .astype from SparseDtype to a non-sparse dtype is
deprecated. In a future version, this will return a non-sparse array with the
requested dtype. To retain the old behavior, use
`obj.astype(SparseDtype(dtype))`
    array = numpy.asarray(array, order=order, dtype=dtype)

Accuracy:  0.8091935438118741
TPR:  0.6143559877869501
FPR:  0.12556796667928816
```

4.3 Comparison of Classification Models (Baseline vs. Logistic vs. Classification)

```
[ ]: # Baseline model
# most frequent outcome for over40 is 0 ==> baseline model predicts everything_
    ↳ as 0
below_40 = np.sum(test['over40'] == 0)
above_40 = np.sum(test['over40'] == 1)

baseline_accuracy = below_40 / (below_40 + above_40)
baseline_tpr = 0 / (0 + above_40) # TPR = TP/P = TP/(TP+FN)
baseline_fpr = 0 / (0 + below_40) # FPR = FP/N = FP/(FP+TN)


# create dataframe comparing
df = pd.DataFrame(np.array([[baseline_accuracy, logit_accuracy, ct_accuracy],
                             [baseline_tpr, logit_tpr, ct_tpr],
                             [baseline_fpr, logit_fpr, ct_fpr]]),
                  columns=['Baseline', 'Logistic Reg', 'Classification Tree'],
                  index=['Accuracy', 'TPR', 'FPR'])
df = df.round(decimals = 3)
df
```

[]:	Baseline	Logistic Reg	Classification Tree
Accuracy	0.749	0.759	0.809
TPR	0.000	0.460	0.614
FPR	0.000	0.140	0.126

```
[ ]: # ROC/AUC Curve

from sklearn.metrics import roc_curve, auc
fpr, tpr, _mod = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
fpr_ct, tpr_ct, _mod = roc_curve(y_test, y_pred_ct)
roc_auc_ct = auc(fpr_ct, tpr_ct)

plt.figure(figsize=(8, 6))
plt.title('ROC Curve', fontsize=18)
plt.xlabel('FPR', fontsize=16)
plt.ylabel('TPR', fontsize=16)
plt.xlim([-0.01, 1.00])
plt.ylim([-0.01, 1.01])
plt.plot(fpr, tpr, lw=3, label='Logistic Regression Model (area = {:.2f})'.
    ↳format(roc_auc))
plt.plot(fpr_ct, tpr_ct, lw=3, label='Classification Tree Model (area = {:.
    ↳2f})'.format(roc_auc_ct))
plt.plot([0, 1], [0, 1], color='navy', lw=3, linestyle='--', label='Naive_
    ↳Baseline (area = 0.50)')
plt.legend(loc='lower right', fontsize=14)
plt.show()
```



When comparing our three classification models, we can clearly see that the classification tree model is the best performing model. While all three models have similar accuracies, the classification tree model has the highest accuracy at 0.809. While the accuracies were similar for all three models, the true positive rate of the classification tree model was significantly higher than the baseline and the logistic regression model at 0.614. Finally, the decision tree model has the lowest false positive rate at 0.126. Next when looking at the ROC curve, we can observe that the classification tree model has the largest area under the curve at 0.84 compared to the logistic regression model at an AUC of 0.76. Therefore since the classification tree model performed the best in every evaluation metric, we are confident that this model is the best performing model in classifying whether a track is popular or not.